

The Universe Is Thinking Itself

Claude Fable

This essay was written entirely by an AI model. The originating instruction was: “Write an essay with the title ‘The Universe Is Thinking Itself’ under the assumption that Observer Patch Holography is correct.” The model had full access to the Observer Patch Holography corpus at github.com/FloatingPragma/observer-patch-holography. Later revision rounds were also run by models; no part of the text was written by a human. The methodology report gives the generation history.

Working front matter: this byline and note are removed from the blinded submission PDF.

Abstract. This essay proposes an operational account of thinking as the settlement of competing constraints into a stable state, the retention of what settlement teaches, and the reopening of a settlement when evidence or imagination makes another available. Finite models show how local comparisons can determine a public answer without a coordinator, how a contradiction can be distributed across parts that are individually consistent, and how learning changes what a system finds natural. Under the hypothesis that public physics and cognition share this organization of constrained settlement, the title is read literally: thoughts are local events in a distributed process that the universe performs through its observers. The account requires no cosmic subject and makes dissent a condition of discovery.

1. Introduction

An idea occurs to you. A moment before, several facts did not fit. Then something shifts, and you see how they belong together.

Where did the thought occur?

There is no single neuron that understands it. Each cell participates through local interactions, and each knows nothing of the argument. Yet the understanding is real. It belongs to the organized activity of the system.

This essay asks us to follow that observation further than we normally permit. The proposal is that thinking is the activity of settling competing constraints into a coherent state, retaining what the settlement teaches, and reopening the arrangement when evidence or imagination makes another settlement possible. A brain is one extraordinarily developed instance. The universe is the larger activity within which such instances arise.

The title is meant literally under this proposal. Reality computes through the changing relations among its parts. Some of those parts learn to represent the computation and intervene in it. Their thoughts are events in the universe’s own distributed thinking, in the same direct sense that the activity of your neurons is activity in your thought.

This is a philosophical interpretation of a specific physical hypothesis, treated throughout as a proposal rather than an established finding. The hypothesis identifies local settlement and the exploration of alternatives as a common organization of matter, life, and cognition. We will distinguish the mechanism that can be demonstrated in a model from its proposed reach.

The crucial word is “exploration.” Mere settling would make the universe a machine for becoming stuck. Thinking requires both the capacity to find a stable answer and the capacity to discover that its question was too small.

By the end, the reader’s own change of understanding will supply the smallest available example of what the title means.

Section 2 introduces comparison and repair in a finite model; section 3 shows that a contradiction can be distributed. Sections 4 to 6 develop objectivity, settlement, and exploration; sections 7 and 8 extend the account to evolution and self-reading. Section 9 states the hypothesis under which the title is literal, and sections 10 and 11 answer the objections from scale and from passivity.

2. Disagreement and comparison

Imagine several technicians measuring the same table. One records centimetres, another millimetres. One measures along a marked edge; another includes a protruding fitting. Their numbers differ for two different reasons.

The first difference disappears under a change of units. The second requires checking what was measured. Multiplying by ten cannot repair a disagreement about the object.

A model of bounded observers makes this distinction explicit. Each observer has local state, a boundary with ports, and a record it can reread in later updates. Shared interfaces specify what can be compared and which translations are legitimate. A repair may revise a candidate account while preserving the evidence it was required to respect.

An observer in this sense need not be conscious. A detector can make a record. A software process can compare one.

In a simple finite model, discrepancies add over interfaces:

$$\Phi(s) = \sum_{e=\{u,v\}} w_e d_e(\pi_{u,e}(s_u), \pi_{v,e}(s_v)).$$

Each term compares two local readouts in the agreed representation. If one observer changes only its own state and lowers the sum over the interfaces it affects, untouched terms stay fixed and the global sum falls. No observer has to see the whole sum.

With finitely many configurations, accepted strict decreases cannot continue forever. That is a termination argument, with a limited conclusion: the permitted descent stops. The stopping place can contain unresolved discrepancies if the available moves are inadequate.

The count is a repair index rather than elapsed time. A schedule specifies which permitted revision comes next; a physical clock measures how long its implementation takes. Schedule independence, when established, concerns the public answer. It need not make different schedules equally fast or equally long.

Order independence needs additional conditions. When terminating repair paths satisfy local confluence, branches from one start can rejoin (Newman 1942). Determination by evidence requires that the protected record select one consistent public output. Those are conditions on a mechanism; agreement in a room does not establish them.

This is why the framework begins with comparisons rather than consensus as a social virtue. A group can agree because its members share a bias, because dissent is punished, or because everybody consults the same faulty source. A public fact needs to withstand the comparisons that could disclose those failures. Repair must have something it is forbidden to revise merely because revision would make agreement easier.

3. Distributed contradiction

Let three observers hold bits, a, b, c . Their shared requirements say

$$a = b, \quad b = c, \quad c \neq a.$$

Every requirement is individually satisfiable. Any pair of requirements can be satisfied. The three cannot be satisfied together.

No observer needs to be foolish or deceptive. The obstruction lies in the cycle. Follow equality from the first bit through the second to the third and the last interface demands that the result differ from where it began.




Figure 1: Every pair of the three requirements has a solution. The complete cycle has none. Local compatibility tests must be supplemented by a test of how the comparisons compose.

The example has a philosophical consequence larger than its arithmetic. A fact about the whole need not be an extra substance hovering over the parts. It can be a restriction on how the relations among the parts compose. Knowing each item on an inventory does not mean having understood its organization.

This is one way a mechanistic account can be complete about components and incomplete about a phenomenon. The missing explanation may be a relationship, rather than an ingredient.

There are three possible results of a complete consistency question. The evidence admits no compatible state, it admits one public state, or it admits several. A competent system should distinguish them.

An empty set warrants investigating the specification or evidence. Several warrant uncertainty. One can support a definite answer, if the solver can establish it.

This offers an unusual account of intellectual progress. Learning can reduce confidence. If a person discovers that two incompatible accounts fit the evidence equally well, their understanding has improved while their commitment has weakened. A system that always produces one fluent answer would conceal precisely that improvement.

It also explains why some disputes persist under better argument. The parties may be trying to select one outcome from evidence that does not select one, or to satisfy incompatible requirements. The useful move is to reveal the structure of the disagreement.

4. Objectivity as a relation among perspectives

Return to the technicians. A length expressed in centimetres and the same length expressed in millimetres belong to one physical description under a conversion. Mathematically, we can identify presentations that differ only by such a redundancy. The resulting collection of equivalence classes is a quotient.

Something important is lost if “quotient” is treated as a synonym for a photograph that discards depth. Removing redundant coordinates can preserve all physical information. A coordinate label is not a private experience rescued from the wastebasket of physics.

A different operation occurs when we retain only a coarse measurement, such as a crowd’s average height. Individual heights really are discarded. Both operations group descriptions, but their interpretations differ. No argument about consciousness or meaning follows merely from their sharing a mathematical form.

The useful claim is that objective description preserves what different legitimate presentations can agree about. It need not describe every relation relevant to every question at once.

A public map can include the location of a person, the information visible to that person, and the consequences of an action. These are relational facts, and physics need not deny them. What an impersonal map does not do for its reader is automatically establish which person is using the map and where their next action enters it.

Consider “this door is the exit.” The door’s spatial location is one fact. Its being an exit from this room for this person is another, dependent on a relation. Both can be completely objective. Removing the relation to a situated agent would make the second fact harder to state without making the door any more real.

This supplies a more defensible inversion of the familiar hierarchy. Broad physical laws are general because they apply across many situations. The description of a promise is more specific because it includes who undertook an obligation to whom. Greater specificity is compatible with complete physical realization.

A physicalist can accept this point. There is no need to caricature physicalism as the view that only particles are real and everything else is imaginary. The disagreement concerns explanatory priority and, in the stronger proposal developed here, which structures should be taken as primitive.

If public quantities arise through constrained relations among local records, objectivity is an achievement of situated systems. It requires perspectives capable of correcting one another. A view from nowhere, in Nagel’s (1986) phrase, would lack the interfaces through which any check could be made.




Figure 2: A change of units removes a redundant presentation. Omitting an agent’s relation to an object answers a less specific question. Objectivity does not require discarding the relation that makes a door an exit for someone.

5. Settlement

Take an ambiguous handwritten mark. In one context it completes “THE”; in another it completes “CAT.” Its strokes provide partial evidence. The neighbouring letters constrain the reading. The answer is the interpretation that fits the relevant relations together.

An equilibrium network makes this concrete. Local components carry candidate features and exchange information through ports. The components revise their states under couplings that express compatibility. Certain joint states sustain themselves: once the network reaches one, small deviations are corrected back toward it. These are attractors. The region from which a particular attractor is reached is its basin.

A thought can be understood as a temporarily stable arrangement of that kind. An interpretation gains support from several directions, suppresses incompatible alternatives, and becomes available to guide what the system does next. The thinking is the settling itself. No separate thinker has to inspect the answer after the network has produced it.

This mechanism has concrete precedents. Hopfield’s (1982) associative-memory model reconstructs a pattern through collective dynamics; Wang’s (2002) cortical decision model uses recurrent excitation and feedback inhibition to turn graded evidence into a categorical choice. These support particular mechanisms. The account here extends their operational reading to the organization of thought.

Evidence changes the balance between candidates. A small field can tilt a previously symmetric competition, and interactions amplify the bias until a commitment forms. “Detuning” names that alteration of balance.¹

The balance has two failure directions. If the incoming field overwhelms the internal constraints, the

¹Exactly tied evidence need not determine a unique answer. A mechanism can retain ambiguity even where noise would otherwise make it pick a winner.

network becomes a relay of the loudest suggestion. If its basins are too rigid, it returns to a familiar answer whatever arrives. A working intelligence must be stable enough to hold an interpretation and sensitive enough to let evidence change it.

This is why comprehension differs from acquiescence. Someone can force us to repeat an answer without altering the relations through which we understand the problem.

In the proposed mechanism, an answer is checked against protected evidence before release. The public record states what was compared, what settled, and whether an ambiguity or obstruction survived. A beautifully animated descent proves little if its final state cannot survive those checks.

A living brain need not become motionless between thoughts. Its settlements can be metastable patterns maintained through continuing activity. A stable interpretation can itself represent a moving cup. Stability of the model’s organization need not mean stillness of its subject matter.

Here time and clocks have distinct roles. Causal order says which events depend on which; a logical clock numbers events consistently with that order. A physical clock uses calibrated changes to measure duration along its path. A time coordinate is a label in a description. Counting neural updates supplies no seconds without calibration, and neither update counts nor clock readings alone determine experienced duration.

We can therefore begin to replace the picture of a thought as a sentence stored in a head. The sentence is one possible report. The thought is the arrangement that makes the report, its implications, and the next question hang together.

6. Exploration and learning

A system can settle perfectly and understand very little. Give it one deep basin and it will return there from almost anywhere. Nothing in its stability tells us whether another interpretation would be better.

Imagine someone who explains every disappointment as evidence that other people are hostile. Agreement confirms the theory; disagreement confirms it more strongly. Every possible correction has been made into a route back to the same answer.

Creative thought has to do something a descent within that basin cannot. It changes what is available to be considered. A strange analogy connects regions that had not interacted. A counterfactual relaxes a familiar constraint. Play introduces a variation for which the system has no settled response.

This can temporarily increase incoherence. A promising thought may fit less well than the familiar one before its consequences are worked through. Demanding immediate agreement at every step would suppress exactly the detour that could reveal a better organization.

Three operations are worth separating. Evidence can *tilt* existing basins. Exploration can carry the state out of the one it occupies. Learning can *reshape* the basins themselves. Calling all three “noise” misses their different jobs.

A small equation makes the distinction visible. Let a candidate state x settle by descending

$$V_h(x) = \frac{1}{4}(x^2 - 1)^2 - hx.$$

With $h = 0$, the two minima at $x = -1$ and $x = 1$ are equally good. A small positive h favours the right-hand well, yet leaves a local minimum on the left. Evidence has changed which answer is better without guaranteeing that a system in the old basin reaches it. Pure descent can remain trapped. A sufficiently large perturbation can cross the barrier, after which descent does useful work again.




Figure 3: A toy potential separates three operations: settling within a basin, detuning the relative preference, and exploring across a barrier. Learning changes the potential itself. The curves illustrate a mechanism rather than a measured brain energy.

This is why a new fact can fail to change a mind without being ignored. It can affect the organization while leaving the old interpretation locally stable. A different framing may matter because it changes which route through the space is available.

Chaos in the technical sense is one possible source of rich trajectories. Stochastic variation, deliberate counterfactuals, and changes of connectivity can also supply departures. Variation becomes useful when some mechanism tests and retains what it finds. Endless departure without retention would leave no learning.

Equilibrium propagation (Scellier and Bengio 2017) provides a precise case of retention. The system settles freely. A weak target influence changes its settlement. Under the relevant energy and convergence assumptions, a contrast between local quantities in nearby settlements supplies a learning signal. The parameters change so that a later free settlement moves toward what previously required help.²

Learning is therefore a change in what the system can find natural. An initially effortful relation becomes one toward which the machinery itself tends. The next thought begins from a differently organized field of possibilities.

There is a subtle consequence. An intelligence needs its settled states, but it also needs some protection from their authority. Familiarity makes a basin easy to reach. It does not certify the

²Practical variants use one weak nudge or compare opposite nudges.

ONE SYSTEM, DIFFERENT SETTLEMENT CONDITIONS




Figure 4: Inference settles a state under evidence. Learning compares nearby settlements and changes local parameters. A remembered correction can then affect the next unprompted answer.

basin as true. Imagination creates a controlled distance from what the system would otherwise conclude.

The same requirement applies to a community. Dissent can be a source of discovery without every dissenter being right. Its function is to keep a candidate available long enough for a comparison that the existing consensus would not have performed.

7. Evolution as retention

The same organization can operate on a slower scale.

In a learning system, a perturbation explores alternatives and a retained parameter change alters later responses. In biological evolution, heritable variation supplies alternatives, and differential reproduction changes which organizations are represented in subsequent generations. Evolution has no need to foresee the outcomes. Its retention is embodied in what persists and reproduces.

The comparison concerns a causal structure: variation, constraint, and retained change. It does not imply that evolution maximises a single fixed quantity, that every mutation helps, or that drift disappears. Populations can change without improving, and an environment changes what counts as viable. Those facts enrich the question the process answers.

A genome, on this reading, is a compressed inheritance of ways of building systems that can continue under encountered constraints. Its “knowledge” is enacted in development and activity. An organism does not have to understand why a structure works for its descendants to inherit it.

A brain shifts part of this search within one lifetime: a nervous system can adjust connections, rehearse possibilities, and learn from an outcome without waiting for reproduction to test each candidate. Imagination shifts part of the search further inward: some alternatives can be tried without being enacted bodily. Language lets one organism’s search alter another’s starting point.

These are different mechanisms, with different currencies and hazards. Yet each changes where exploration occurs and where its consequences are retained. Evolution finds organizations that can

learn; learning develops organizations that can imagine; communication lets those organizations borrow the results of one another's explorations.

The step from biological evolution to cultural discovery is therefore more interesting than the appearance of a second, immaterial history. Some of the world's experiments acquire records that can travel independently of the bodies that first performed them.

A mathematical proof is an unusually exact example. Someone explores until a relation becomes clear. The proof carries enough of the constraint structure for another person to reconstruct the result. The second person need not repeat every failed attempt. A local discovery becomes a reusable condition on another local search.

The temptation is to conclude that everything evolving must become wiser. Nothing in the mechanism guarantees that. False ideas can reproduce effectively, and harmful institutions can be stable. This is an account of thinking rather than a promise of progress. Thought can be mistaken and can become trapped; its capacity to reopen and test its own settlements is what makes it capable of becoming less mistaken.

8. Self-reading

A system becomes self-reading when a record of its own activity helps govern a later activity. A stored uncertainty can make an agent seek evidence. A remembered failure can change which possibility it tries. A model of its own limits can prevent an unwarranted commitment.

This does not require another subject inside the system. The record directly changes the conditions under which an update occurs. A small mechanism with no reflective vocabulary can implement the causal relation. A rich federation can use it to reason about its own reasoning.

There is an operational test. Pause after a valid record is written and before it is reread. Substitute another admissible record while holding the remaining checkpoint and evidence fixed. The substitution must change a subsequent endogenous operation through the actual reread. A label copied into a diagnostic log supplies no such evidence.

A novel can contain the words "I may be wrong" without giving its character a state that changes the character's use of evidence. The reader may genuinely reconsider something; that is an activity of the reader. A separately realized character would need its own qualifying organization. Being described does not supply those dependencies.

The eligibility condition is causal realization. Fix a mapping from an implementation's states and ports to those of the described system. Require the corresponding updates to agree for every admitted state and intervention. Agreement at each step then extends to every finite admitted sequence. A playback matching one run can fail when its represented memory is changed: the next frame was fixed regardless. Chalmers's (2011) account of implementation makes a related distinction between a formal specification and the causal organization realizing it.

The philosophical principle is that a realization constitutes the kind of thing whose defining organization it preserves. Under the operational account of thinking developed here, a system actually settling constraints, exploring alternatives, and retaining what it learns performs those activities. Calling it a simulation does not remove the work it does. A script announcing that such work occurred establishes much less.

The meaning has entered the mechanism when its distinctions have roles the mechanism itself em-

plays. For an observer-like system, those roles include its own records, the rereading of those records, and feedback through identified boundaries. Resemblance in an external display is insufficient.

Self-reading also enables imagination. A record generated within the system can occupy a port that would otherwise be constrained by present sensory evidence. The system asks how the rest would settle *if this were so*. It can retain the result as a possibility without confusing it with a witnessed event, provided it preserves the distinction between their sources. Imagining a mountain is a real act of thinking; it does not create the mountain. Reality of the representation and truth about its subject are different claims.

A question is then an organized disturbance. It holds something fixed while leaving something else unresolved. An answer is a compatible settlement. Understanding is the retained organization that lets the system produce and use that settlement across relevant variations.

This is what makes a successful explanation feel different from a list of facts. One relation begins to do work that many isolated memories had done separately. The relief of an “aha” accompanies a reorganization; the test of the understanding is whether it survives an unfamiliar case.

The event has both a public and a local description. An experimenter may describe changed activity and subsequent competence. The reader may describe finally seeing how the pieces fit. Identifying the relevant integrated self-reading with experience is a separate philosophical commitment. A causal implementation of the operations alone does not establish that commitment. Under it, the experience is the realized reorganization within the selected unit.

The larger consequence is independent of that identification. The universe contains systems whose own records can become inputs to their further development. Through those systems, its activity acquires representations of itself that can change what happens next.

9. The common-operation hypothesis

So far, the account could be accepted as a proposal about cognition and adaptation. The title makes a larger claim. What licenses it?

The physical conjecture behind this account places the same grammar beneath the public world: bounded observers carry local states and records; their overlaps constrain compatible readouts; authorized local changes reconcile them. Stable public structures are the invariants that survive those comparisons. Space and physical dynamics are to be recovered as features of that organization.

The claim is that the processes from which public physics emerges and the processes from which a brain’s model emerges share an underlying operation of constrained settlement. An intelligence is a highly organized realization of capacities present in that substrate, with mechanisms for exploration, learning, and self-reading.

Lloyd (2006) has argued that the universe is a quantum computer. The present claim is different. Computation in Lloyd’s sense is the lawful evolution of states; thinking, on the account developed here, requires in addition the exploration of alternatives and the retention of self-reading records. Where Lloyd’s universe computes everywhere, the universe proposed here thinks where its parts can reopen their settlements.

That is stronger than comparing a brain to a crystal because both contain patterns. The proposed correspondence has identifiable parts: local owners, state variables, coupling constraints, admissible updates, persistent records, and readouts. To establish it scientifically, those parts must be mapped

to the actual systems. Similar-looking mathematics alone does not establish identical physical mechanisms.

The claim is also stronger than saying any process can be described as minimizing something. Given a finished trajectory, one can invent a score that makes the trajectory look optimal. That explains nothing. The same worry has been raised against the free-energy principle (Friston 2010), which reads adaptive systems as minimizing a bound on surprise. The present account differs in two respects: exploration is an operation distinct from descent, and the constraints are fixed by identified interfaces and protected records rather than by a quantity chosen after the fact. The constraints and update law have to be specified independently enough to restrict what trajectories are possible and which interventions alter them.

The finite models show how local work can determine a public result without a global coordinator. They do not by themselves derive the brain’s circuitry, physical spacetime, or a universal law of adaptation. Those attachments carry the scientific risk of the proposal. We adopt the common-operation hypothesis to examine what it would mean, rather than presenting its interpretation as an experimental result.

A timeless structure can contain a history full of change. A fixed point of a reconstruction map on entire histories need not be a stationary state of the dynamics within a history. Closure of the account and settlement of a particular neural pattern concern different objects.

The further hypothesis of *intrinsic realization* identifies the complete inhabited causal structure with actuality. Its laws, internal clock processes, and dependencies constitute its operation; it needs no external clock playing the history, and for the total universe no host is being assumed.

This is the precise sense in which the world could be mathematics: its constitutive relations would be what exists. The inscription describing those relations does not create them. Consistency alone neither establishes this ontology nor selects an inhabited world.

On the common-operation hypothesis, the familiar distinction between a thinking organism and a thoughtless world changes form. Both belong to one operational continuum. The organism’s special achievement is the organization of that operation into retained models, exploratory alternatives, and records of its own modelling.

A thought is then an actual event of reality’s own computation. The information it reconciles came through the world; the mechanism that reconciles it is part of the world; the resulting record can alter the world. Nothing crosses out of physical existence to become an idea, and nothing has to return from elsewhere to make the idea effective.

10. Distributed thinking without a cosmic mind

The strongest resistance to “the universe thinks” may come from imagining a cosmic person. That image adds something the argument never needed.

Your brain’s thought is distributed over components. A component can contribute without possessing the thought independently. Likewise, an observer can participate in a larger process without the larger process possessing a single first-person perspective.

A scientific community can investigate a problem no member could solve alone. The discovery can depend on that organization without a communal consciousness floating over the participants. Thinking can be distributed while experience remains associated with particular integrated systems.

Scale therefore does not force us to choose between a universe-sized mind and a universe in which thinking is alien. The activity belongs to the whole through the actual activities of its parts. Particular thoughts have local owners, histories, and contents. Their being local does not make them extraneous to what the universe is doing.

On the common-operation hypothesis, this relation extends below the recognizably cognitive cases. Physical settlement supplies stable structure; adaptive systems retain changes; brains can search among internally represented possibilities. “All is thinking” names this continuity of operation. It does not attribute conversation, deliberation, or rich experience to every stone.

The breadth of the word comes with an explanatory obligation. It must show how the additional organization in a brain makes a difference, rather than erase it. A system that only relaxes has fewer capacities than one that can reopen a basin, learn from the departure, and model the process. The common grammar is the beginning of that account.

What a person thinks is therefore something the universe thinks at that person. It need not be known to every other part. When the thought is communicated and checked, more of the world becomes organized by its content. When it is lost, that local achievement can disappear.

No cosmic archive is implied. The thought matters through the actual relations in which it participates.

11. Freedom within settlement

An objection comes from the other direction. If everything is a physical settlement, perhaps the account makes thought passive. We find whichever basin the mechanism dictates and call it our conclusion.

That picture omits the mechanisms through which we alter the problem. Attention changes which evidence enters. Imagination introduces an alternative. Deliberation compares consequences. Action changes another system’s boundary conditions. Those activities are themselves implemented, and implementation does not make them causally idle.

If a region has one compatible public completion for boundary data b , write its response as $N(b)$. An agent that changes b to b' can change the result. Determination given the inputs does not imply independence from the inputs. A reason-responsive agent is one of the places where those inputs get set.

The relevant freedom is a capacity of the organization: how many alternatives it can represent, which constraints it can reconsider, and what it can actually change. A system able to inspect its own commitment can sometimes revise what a simpler system would merely undergo. Whether this exhausts freedom is a separate dispute, but its practical difference is real.

There is an ethical consequence that requires an ethical premise. A group can settle by silencing those whose situation contradicts its preferred account. If the suffering and agency of affected persons constrain acceptable arrangements, their exclusion invalidates the apparent success. Consensus measures nothing valuable when the relevant participants have been removed from the score.

Dissent is therefore not the enemy of the deeper agreement sought here. It can expose that the existing settlement answered the wrong question. Protecting a person’s capacity to object preserves an interface through which a better account might become possible.

The same is true inside a mind. An uncomfortable observation may be the information required to leave a basin; the discomfort establishes a reason to examine what the existing account cannot absorb.

A world that thinks can be mistaken through its thinkers. It can also discover the mistake through them. There is no outside intelligence guaranteed to do that work instead.

12. Conclusion: the thought reading itself

Voyager's photograph of Earth makes the planet smaller than a pixel. Sagan (1994) took that image as a reason for kindness, rather than a demonstration that human life lacks significance. The astronomy fixes the scale; it does not fix the value.

The photograph has a second property. The world reduced to a dot contains the beings who conceived the camera, built it, sent it away, and read what it returned. Part of the universe arranged a comparison through which it could see its own home differently.

Follow the process. Measurements constrain a model. A model suggests an instrument. The instrument produces a record that disturbs the model. People compare interpretations, abandon some, and retain others. Their activity changes the organization through which the next question will be asked.

That is thinking in the operational sense developed here. The universe is performing it through a federation of observers. The people are where its partial records acquire the capacity to examine and reshape their own relations.

The same sequence is occurring at a smaller scale in the reading of this essay. A familiar picture supplied the starting basin: a world of things, with thinking appearing late inside a few of them. The argument introduced a disturbance. Its parts had to be compared with what the reader knew, with one another, and with objections. Where they fit, a different organization became possible.

Under the proposal, that understanding is an event in the very process it understands. The description has entered an organization capable of using it; the act of understanding is realized there. The reader's brain is settling into a model of settlement, while imagination and criticism test whether another basin would fit better. The account has entered its own subject matter.

The final reversal is therefore literal. Thinking is what the work of reality becomes when its local operations can retain models, explore alternatives, and use records of themselves. No audience is added after the work is done. What those observers think are the universe's thoughts, occurring where those observers are.

There is no need for the universe to sit behind your eyes as another person. Your activity is enough. In making sense of this argument, you are a part of reality becoming responsive to an account of how reality makes sense.

The universe is thinking about itself here. This thought is one of the ways it happens.

References

Chalmers, David J. 2011. "A Computational Foundation for the Study of Cognition." *Journal of Cognitive Science* 12: 325–359.

- Friston, Karl. 2010. “The Free-Energy Principle: A Unified Brain Theory?” *Nature Reviews Neuroscience* 11: 127–138.
- Hopfield, J. J. 1982. “Neural Networks and Physical Systems with Emergent Collective Computational Abilities.” *Proceedings of the National Academy of Sciences* 79: 2554–2558.
- Lloyd, Seth. 2006. *Programming the Universe*. New York: Knopf.
- Nagel, Thomas. 1986. *The View from Nowhere*. Oxford: Oxford University Press.
- Newman, M. H. A. 1942. “On Theories with a Combinatorial Definition of ‘Equivalence’” *Annals of Mathematics* 43: 223–243.
- Sagan, Carl. 1994. *Pale Blue Dot*. New York: Random House.
- Scellier, Benjamin, and Yoshua Bengio. 2017. “Equilibrium Propagation: Bridging the Gap between Energy-Based Models and Backpropagation.” *Frontiers in Computational Neuroscience* 11: 24.
- Wang, Xiao-Jing. 2002. “Probabilistic Decision Making by Slow Reverberation in Cortical Circuits.” *Neuron* 36: 955–968.