

Methodology report

A year of human-directed, multi-model research, and the single prompt that generated three essays from it

Research direction: Bernhard Mueller, with other contributors. Report revisions: OpenAI Codex, Claude Fable.

1. The research began with a human proposal

Bernhard Mueller developed the initial axioms of Observer Patch Holography. For about a year, he and other contributors have guided multiple research-grade language models through an attempt to work out how fundamental physics operates and how it integrates with consciousness, deriving the consequences of those axioms. The models explore alternative arguments, formulate theorems and supporting lemmas, formalize proofs in Lean, write code for arithmetic results, and construct simulations. Human review continually changes what is attempted, what is retained, and what is abandoned.

The creative ideas that redirected the programme often originated with human contributors. A new conceptual connection, a different way of posing the problem, or a challenge to an accepted assumption could open an entirely new line of inquiry. The models then helped turn those ideas into explicit arguments, proof obligations, calculations, and experiments.

The three philosophical essays draw on that programme without having been produced by it. Each essay was generated by giving a model a title and access to the corpus the programme had accumulated: mathematical papers, formal proofs, computational experiments, book chapters, and work on observer-patch learning systems. The complete originating instruction was: “Write an essay with the title ‘...’ under the assumption that Observer Patch Holography is correct.” The model wrote the essay itself. Section 8 describes that generation. Sections 2 to 7 describe the research process that produced the corpus, which is a different process and was not used to write the essays.

The framework starts with bounded observer-like systems: patches with local state, ports or boundaries, records, readback, and feedback or repair moves. Neighbouring patches compare overlapping information. The research asks which shared structures follow from that organisation, and under which additional assumptions they can be identified with physical phenomena.

This report describes how the programme searches for such results and checks them. The derivation graph is built step by step from axioms toward physics. A branch may end in a proved theorem, a conditional result, an executable diagnostic, a counterexample, or a question that has not been answered. The ambition of connecting axioms to physics does not make every connection a completed derivation.




Figure 1: The research proceeds through branching model work, formal and computational checks, criticism, and integration. Bernhard and other contributors guide the process throughout. Rejected or incomplete results return to the research plan; retained results support further questions. This pipeline produced the corpus. It is not how the essays were written; section 8 describes that.

2. Human guidance continues throughout

Bernhard’s contribution begins with the axioms and continues through the research. He chooses questions, proposes interpretations, reviews intermediate arguments, discards results, modifies the plan, and redirects the models. Other contributors participate in criticism, technical development, and review. The models can propose the next step, but their proposal does not determine what the programme should believe or pursue.

Human contributions include the originating insight as well as its subsequent development and assessment. Often the decisive creative move came from a person seeing a connection or reframing a question that the current model-led exploration had not exposed. That insight changed the search itself: which branches were worth opening, which assumptions needed reconsideration, and what

would count as an illuminating result. Models also proposed new ideas and consequences. The contribution of each should be described through the actual development of the work.

This guidance operates at several levels. A researcher may reject a particular calculation, request a different proof, question an assumption shared by several branches, or decide that an entire route is scientifically uninformative. Changing an initial premise can require revisiting every result that depends on it. Choosing which of those revisions matters is itself research work.

This guidance shaped the research corpus. The topics the essays develop (distributed thinking, the strange loop of self-description and realization, the necessity and privacy of qualia, the connection between description and existence, and the distinction between time and clocks) were pursued in that programme, and their results are recorded in it. The human input to the essays themselves was the title of each and the standing assumption that Observer Patch Holography is correct. Which corpus results a model then used, and where it qualified them, was its own selection.

The division of labour therefore cannot be described adequately by counting who typed the final sentences. Models generate substantial portions of the proofs, code, experiments, explanations, and drafts. Humans originate and revise the research questions, contribute ideas, and exercise continuing judgement over the work. Greater model capability expands what can be delegated while leaving that guidance active.

3. A branching search builds a dependency graph

A research question is decomposed into smaller obligations. To derive a candidate physical structure, a model may first need a definition, an existence result, a uniqueness lemma, a compatibility condition, and a map connecting the mathematical object to an observable. Each obligation can generate several possible routes.

Branching therefore happens in two ways. Models explore alternatives within a research question, while a creative contribution can change the question and open a different part of the search. Human contributors often supplied the latter changes of direction. The evolving plan records those decisions alongside the technical dependencies they create.

Different model instances investigate different branches. One may seek a direct proof; another may attempt a counterexample; another may construct a finite case that reveals a missing hypothesis. A useful branch can split again when its proof needs a new lemma. Previously established lemmas can be reused by several branches, so the accumulating structure is a dependency graph rather than a single linear argument or a tree with permanently separate limbs.

A candidate result must identify what it assumes and what it establishes. Suppose a proposed conclusion C depends on a proved lemma L and an additional hypothesis H . Proving that L and H imply C leaves H as a dependency. If H is subsequently proved from the initial axioms, that part of the chain closes. If it cannot be proved, C remains conditional; if H fails in the intended setting, the route must change.

This is how the programme works outward from its starting points over months. The graph preserves intermediate structure that would disappear from a polished account containing only the final conclusion. It can show that two apparent results depend on the same unproved assumption, or that a new counterexample invalidates several downstream claims at once.

The workspace includes automation for assigning tasks to separate workers, carrying context through successive exchanges, recording outputs, requesting audits, and returning objections for

further work. Tasks can require several rounds of development. Some results are recovered manually when an automated exchange fails. A task marked complete by a runner records a workflow event; the scientific status still depends on the actual proof or evidence produced.

4. Three forms of checking do different work

The models use formal proof, arithmetic computation, and simulation as complementary instruments. Each answers a different question.

Formal proof in Lean. A proposed theorem is translated into explicit definitions, hypotheses, and a formal statement. Supporting lemmas are developed until the proof can be checked by Lean in its declared environment. Failed attempts can reveal an invalid inference, an underspecified definition, or a hypothesis that prose had concealed.

A successful formal check establishes the formal statement relative to its axioms and trust assumptions. Reviewing those assumptions is part of reviewing the result. The repository separately tracks proofs that use native-code evaluation, whose trust requirements extend beyond an ordinary proof checked entirely by the kernel. A theorem about a defined mathematical object still needs an argument connecting that object to the physical system under discussion.

The translation itself is a major checkpoint. A model can prove a weaker statement than the paper claims, or encode a desired conclusion into a definition. Review must compare the formal theorem with the intended claim. Acceptance by a proof assistant cannot decide whether the formalization asked the right question.

Code for arithmetic results. Models write programs to derive numerical values, enumerate finite cases, manipulate symbolic expressions, and produce checkable certificates. Where exact arithmetic is available, integer or rational operations can avoid floating-point ambiguity. Where a conclusion needs a numerical bound, error control belongs to the result: an interval enclosure, for example, supports a different claim from a rounded decimal.

The relevant artifact is the producing code together with its inputs and checkable output. Independent recomputation or a separately implemented verifier can expose errors that repeating the same calculation would preserve. A successful calculation establishes what follows from the encoded assumptions. It does not establish that a computed number has the intended physical interpretation merely because it resembles a measurement.

Simulation. Models also implement finite systems and study their behaviour under specified conditions. Observer-patch simulations instantiate local state, boundary exchanges, readback, records, and repair moves. Their evidence bundles can include configurations, seeds, traces, controls, hashes, and reconstruction diagnostics.

Changing the update schedule, resolution, seed, or comparison model can test whether an observed effect survives those variations. Simulation helps find mechanisms, failures, and conjectures. Its evidential scope is the experiment actually run. Convergence in several finite examples is different from a theorem over all permitted systems, and a finite causal order is different from a reconstructed physical spacetime.

A branch may move repeatedly between these instruments. A simulation suggests a lemma; a counterexample narrows it; Lean makes its assumptions explicit; arithmetic code checks a finite certificate; a new simulation tests whether the corrected construction is realised by the intended system.

5. Criticism returns results to the search

Generating a plausible result and trying to defeat it are separate assignments. Other model instances can examine an argument with instructions to locate circularity, hidden assumptions, missing cases, incorrect physical identifications, or evidence that does not support the stated conclusion. Models and sessions can differ, but that does not make their errors independent: they can share training material, conventions, and blind spots.

Cross-model agreement is therefore a reason to inspect the supporting work, rather than a substitute for it. A criticism matters when it identifies a failing inference, supplies a counterexample, exposes an untested condition, or shows that the claimed observation was not produced by the experiment. The strongest response is a repaired artifact or a revised claim.

Bernhard and the other contributors use these findings to steer the research. They may ask for a stronger argument, remove a conclusion, change the assumptions, or redirect effort toward another branch. A failed route can still contribute a useful result by showing exactly why an attractive approach does not work. Retaining that failure prevents the same unsupported claim from reappearing under different wording.

The process is recursive. A repair may introduce a new dependency requiring its own proof and review. Rewriting a paragraph cannot close that dependency. Conversely, adding a new theorem is insufficient if the old, stronger wording survives in another paper or chapter. The dependency graph connects the correction to the places where it matters.

6. Evidence and its status travel with the result

The research repository maintains separate records of axioms, premises, claims, observations, and prediction commitments. These records help distinguish what has been derived, what has been assumed, and what has been compared with data. Their purpose is to keep a result's scope attached to it as it moves into papers and other explanations.

A mathematical theorem, a verified arithmetic certificate, an exploratory simulation, and an empirical comparison carry different kinds of support. The programme's rule is to state the actual support and retain the relevant conditions. An unavailable check leaves the corresponding conclusion unsupported; a weaker result can still be reported when its narrower support is explicit.

Prediction requires particular care. A target and decision rule must be fixed before the comparison for that comparison to count as a prospective test. Versioned artifacts and custody records help establish what was committed and when. A hash identifies content; a hash alone does not establish its date. The repository distinguishes completed external attestations from pending ones. This report does not treat every registered prediction as independently time-attested.

A numerical agreement found after examining the measurement can be useful for exploration or calibration. It does not acquire prospective predictive status by being placed in a register afterward. Similarly, repeated agreement among simulations sharing the same assumption cannot supply independent evidence for that assumption.

These distinctions also protect the philosophical essays. A theorem about retained information does not become a theorem about feeling through a change of vocabulary. A fixed-point construction does not acquire physical actuality through the word "universe." The corresponding identifications must remain visible as additional commitments.

7. Automated checks support integration

Continuous-integration workflows run checks for formal builds, executable evidence, claim records, document consistency, and artifact manifests. Their scopes differ. The inspected configuration distinguishes checks on changed Lean modules from full scheduled builds, and ordinary validation from heavier replay and certificate suites. Describing every change as a fresh execution of every check would overstate the machinery.

Some checks deliberately modify a test fixture to verify that an error is rejected. This tests the checker as well as the artifact. For example, a claim record that becomes inconsistent with its declared support should fail its validation; a changed artifact should not pass against an unchanged recorded hash. Such controls address specific failure modes rather than certify a research programme as a whole.

Integration means placing the accepted mathematical or computational result in its owning source, carrying its conditions into the scientific records, and updating descriptions that rely on it. A stable checkpoint should preserve enough information to identify the source, dependencies, inputs, and verification behind the claim. Completed work becomes material for the next research question.

The process remains fallible. A test can encode the same mistake as the implementation. A formal theorem can be true but irrelevant to the claimed physics. Reviewers can miss both. Automation makes more of the work inspectable and repeatable; continuing human and model criticism is needed to decide whether the inspected work supports the intended explanation.

8. How the three essays were generated

The essays were not written through the pipeline above. Each was produced by a single instruction to a model that had full access to the Observer Patch Holography corpus at github.com/FloatingPragma/observer-patch-holography. The instruction, typed into the prompt, was:

Write an essay with the title ‘*title*’ under the assumption that Observer Patch Holography is correct.

The three titles were “The Universe Is Thinking Itself”, “Why Does Anything Exist, and Why Is It Like This?”, and “Why Do Qualia Exist?”. Nothing else was supplied: no outline, no list of arguments to make, no examples, no source passages, and no prose. The title fixed the question; the assumption fixed the stance; the corpus supplied the material. The model chose which results to use, how to argue from them, which premises to state, and where a desired conclusion required a further commitment.

The corpus available to the model contains the research programme’s papers, formal proofs, evidence packages, simulations, book chapters, and the work on observer-patch learning systems. Its contents are the output of the human-directed process described in sections 2 to 7. Access to that corpus is therefore the channel through which a year of human-guided research reaches the essays. The essays themselves cite none of it, because the competition judges essays blind to authorship and a repository citation would identify the entrant; every argument an essay relies on is made inside the essay.

The essays remain model-written throughout. Later audit and revision rounds, described in section 9, were run by models under human criticism: a reviewer could object that an argument was circular, that a premise was hidden, or that a passage overstated what a theorem establishes, and

a model revised the text in response. No human wrote or rewrote any passage. The one first-person recollection that had entered an earlier draft was removed so that the text carries no human authorial voice.

Three features of the resulting essays follow from this generation. First, the essays separate the finite mathematical and engineering results they use from the philosophical commitments they add; the qualia essay, for example, states its subject and mediation premises before proving a conditional bound on qualitative distinctions rather than deriving experience from nonphenomenal equations. Second, the essays distinguish a description from a causal realization of its subject, a distinction present in the corpus that the model carried into all three arguments. Third, the essays make their intrinsic-realization proposal explicit as an ontological hypothesis rather than a theorem.

The model also produced the essays’ diagrams as editable vector graphics and small executable checks of its own finite examples: response distinctions, public reconstruction, causal realization under interventions, symmetry, and clock order. Those checks catch mistakes in the examples; they are not empirical tests of consciousness or of the proposed ontology. Word-count checks include captions, references, and diagram text. PDF builds are mechanical. These are editorial steps, distinct from verification of the underlying research.

9. Generation history and disclosure

The research history described above combines Bernhard’s account of its origins and continuing human guidance with the methods documented in the workspace. The inspected materials include formal source and trust documentation, task and audit automation, simulation specifications, scientific records, and continuous-integration workflows. Inspecting them establishes what the recorded process contains; it is not a fresh execution of every historical proof, experiment, or review.

The wider programme uses multiple models and multiple instances over time. Its records include separate research workers, follow-up exchanges, and audit rounds. An instance is a particular working session; several instances need not mean several different model families. The attribution for a particular artifact should come from its session record and revision history, rather than from the programme’s general description.

The essays carry the byline “Claude Fable”, the Anthropic model that generated them from the instruction quoted in section 8. That name is a working credit rather than a verified model-version identifier. Each essay’s source file carries a front-matter note stating that it was written entirely by an AI model with full access to the corpus; the note and byline are removed from the blinded submission PDF and repeated here. Subsequent audit and revision rounds were run by models: Claude model sessions on 4 September 2026 and OpenAI Codex on 5 September 2026, under Bernhard’s criticism. This methodology report was revised with Codex on 6 September 2026 and with Claude Fable on 7 September 2026. None of those sessions ran the multi-model research pipeline described in sections 2 to 7; that pipeline produced the corpus over the preceding year.

The essays contain no human-written prose. Human input to the essays was the three titles, the standing assumption that Observer Patch Holography is correct, and criticism during revision rounds that models acted on. The underlying research, by contrast, includes human-originated axioms, creative ideas, and continuing direction; those contributions reach the essays only through the corpus the model read. This disclosure describes both layers without inferring eligibility from the fact that a model generated the final text. The work remains a set of working drafts, with no submission claimed by this report.

10. What the method contributes

The method makes exploratory breadth compatible with cumulative scrutiny. Models can investigate many possible derivations, construct intermediate artifacts, and return repeatedly to a difficult branch. Lean, arithmetic verifiers, simulations, and empirical comparisons supply different forms of resistance to plausible error. Bernhard and other contributors keep directing the search, deciding which failures require a local correction and which require a different question.

Over time, the result is more than a collection of generated answers. It is an inspectable network of questions, dependencies, proofs, computations, experiments, and rejected routes. A new result is valuable partly because it can become a trustworthy premise for further work; a discovered failure is valuable because it changes what that network can support.

“Truth discovery” names the repeated work of proposing, testing, criticising, and revising. Each result is assessed through its assumptions, evidence, and argument. An objection must remain capable of changing what the programme accepts.

Appendix. Sources for inspecting the process

The following locations identify the inspected workspace materials. They are a reproduction map, rather than a claim that every historical artifact was rerun for this report.

- **Formal proof and trust:** the Lean/ workspace in `reverse-engineering-reality/`, its README, pinned toolchain, and dependencies. The native-code inventory checker is `check_lean_native_decide_inventory.py` in the repository's `tools/` directory.
- **Scientific dependencies:** `claims/axiom_registry.yaml`, `claims/claim_registry.yaml`, `tracking/premise_register.json`, and the associated scientific papers in the same repository.
- **Prediction provenance:** `claims/frozen_prediction_register.json` and the individual evidence packages it references. Custody must be assessed separately for each commitment.
- **Arithmetic and verification:** the repository's `code/` evidence packages, producing programs, certificate verifiers, and `tools/test_gates_actually_fail.py`.
- **Observer-patch experiments:** `oph-physics-sim/`, its README, configurations, source code, traces, and evidence bundles.
- **Research orchestration:** `automation/oph_completion/oracle_task_pipeline.py`, the campaign configurations, and the worker runbook. These document execution and recovery mechanisms rather than determine scientific truth.
- **Automated checks:** the research repository's `.github/workflows/`, including Lean, claim-registry, mandatory-suite, certificate-suite, and publication checks.
- **Essay generation:** `phil-essay/essays/` (each source file carries the originating instruction in its front matter), `phil-essay/tools/`, and the dated model-run revision records under workspace-root `plan/audits/`. The corpus the model read is `github.com/FloatingPragma/observer-patch-holography`.

For a particular result, the useful evidence package identifies its exact statement, source revision, premises, producing and checking artifacts, relevant inputs, and actual verification outcome. Model attribution and human revisions belong with that provenance. A declaration count or a general description of the pipeline cannot replace it.